solution

October 9, 2024

```python
[11]: # Import necessary libraries
      import os
      import numpy as np
      import pandas as pd
      import matplotlib.pyplot as plt
      import cv2
      from sklearn.model_selection import train_test_split
      from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,␣
       ↪Dropout
      from tensorflow.keras.preprocessing.image import ImageDataGenerator
      from tensorflow.keras.callbacks import EarlyStopping
      import glob
```

```python
[12]: # Set constants
      IMG_SIZE = (128, 128)  # Image size for resizing
      BATCH_SIZE = 32
      EPOCHS = 10  # Adjust as needed

      # Define file paths
      TRAIN_LABELS_CSV = 'train_labels.csv'
      TRAIN_DIR = 'train'
      TEST_DIR = 'test'
      SUBMISSION_CSV = 'submission.csv'
```

```python
[13]: # Load labels from CSV
      train_labels = pd.read_csv(TRAIN_LABELS_CSV)
      print(train_labels.head())
```

```
                                         id  label
     0  f38a6374c348f90b587e046aac6079959adf3835      0
     1  c18f2d887b7ae4f6742ee445113fa1aef383ed77      1
     2  755db6279dae599ebb4d39a9123cce439965282d      0
     3  bc3f0c64fb968ff4a8bd33af6971ecae77c75e08      0
     4  068aba587a4950175d04c680d38943fd488d6a9d      0
```

```
[14]:  # Exploratory Data Analysis (EDA)
       # Check the distribution of classes
       label_counts = train_labels['label'].value_counts()
       plt.figure(figsize=(8, 6))
       label_counts.plot(kind='bar')
       plt.title('Distribution of Labels in the Training Set')
       plt.xlabel('Label (0: Non-Metastatic, 1: Metastatic)')
       plt.ylabel('Count')
       plt.xticks(rotation=0)
       plt.show()
```

Distribution of Labels in the Training Set


```
[15]:  # Load images in batches
       def load_image(image_path):
           img = cv2.imread(image_path, cv2.IMREAD_COLOR)
           if img is not None:
               img = cv2.resize(img, IMG_SIZE)
               return img
           return None

       def generate_data(batch_size):
```

```python
        while True:
            for start in range(0, len(train_labels), batch_size):
                end = min(start + batch_size, len(train_labels))
                batch_images = []
                batch_labels = []
                for i in range(start, end):
                    img_path = os.path.join(TRAIN_DIR, train_labels['id'].iloc[i] +
 '.tif')
                    img = load_image(img_path)
                    if img is not None:
                        batch_images.append(img)
                        batch_labels.append(train_labels['label'].iloc[i])
            yield np.array(batch_images), np.array(batch_labels)
```

```python
[16]: # Create a CNN model
      def create_model():
          model = Sequential()
          model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(IMG_SIZE[0],
       IMG_SIZE[1], 3)))
          model.add(MaxPooling2D(pool_size=(2, 2)))
          model.add(Conv2D(64, (3, 3), activation='relu'))
          model.add(MaxPooling2D(pool_size=(2, 2)))
          model.add(Conv2D(128, (3, 3), activation='relu'))
          model.add(MaxPooling2D(pool_size=(2, 2)))
          model.add(Flatten())
          model.add(Dense(128, activation='relu'))
          model.add(Dropout(0.5))
          model.add(Dense(1, activation='sigmoid'))  # Binary classification
          model.compile(optimizer='adam', loss='binary_crossentropy',
       metrics=['accuracy'])
          return model
```

```python
[17]: # Prepare training and validation datasets
      X_train, X_val, y_train, y_val = train_test_split(train_labels['id'],
       train_labels['label'], test_size=0.2, random_state=42)

      # Initialize the model
      model = create_model()
      model
```

```
[17]: <keras.engine.sequential.Sequential at 0x7feee3f40130>
```

```python
[18]: # Train the model
      # Note: Training done outside of notebook environment.
      early_stopping = EarlyStopping(monitor='val_loss', patience=3)
      train_gen = generate_data(BATCH_SIZE)
```

```
model.fit(train_gen,
          steps_per_epoch=len(X_train) // BATCH_SIZE,
          validation_data=generate_data(BATCH_SIZE),
          validation_steps=len(X_val) // BATCH_SIZE,
          epochs=EPOCHS,
          callbacks=[early_stopping])
```

2024-10-09 10:49:16.155514: I
tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:185] None of the MLIR
Optimization Passes are enabled (registered 2)

Epoch 1/10

2024-10-09 10:49:17.353270: W
tensorflow/core/framework/cpu_allocator_impl.cc:80] Allocation of 65028096
exceeds 10% of free system memory.
2024-10-09 10:49:17.709597: W
tensorflow/core/framework/cpu_allocator_impl.cc:80] Allocation of 30482432
exceeds 10% of free system memory.
2024-10-09 10:49:18.037437: W
tensorflow/core/framework/cpu_allocator_impl.cc:80] Allocation of 23482368
exceeds 10% of free system memory.
2024-10-09 10:49:18.152410: W
tensorflow/core/framework/cpu_allocator_impl.cc:80] Allocation of 30482432
exceeds 10% of free system memory.
2024-10-09 10:49:18.192791: W
tensorflow/core/framework/cpu_allocator_impl.cc:80] Allocation of 25719552
exceeds 10% of free system memory.

  91/5500 […] - ETA: 48:00 - loss: 6.9329 -
accuracy: 0.6078
```
---------------------------------------------------------------------------
KeyboardInterrupt                         Traceback (most recent call last)
Cell In[18], line 5
      2 early_stopping = EarlyStopping(monitor='val_loss', patience=3)
      3 train_gen = generate_data(BATCH_SIZE)
----> 5 model.fit(train_gen,
      6           steps_per_epoch=len(X_train) // BATCH_SIZE,
      7           validation_data=generate_data(BATCH_SIZE),
      8           validation_steps=len(X_val) // BATCH_SIZE,
      9           epochs=EPOCHS,
     10           callbacks=[early_stopping])

File ~/.local/lib/python3.8/site-packages/keras/engine/training.py:1184, in
    Model.fit(self, x, y, batch_size, epochs, verbose, callbacks,
    validation_split, validation_data, shuffle, class_weight, sample_weight,
    initial_epoch, steps_per_epoch, validation_steps, validation_batch_size,
    validation_freq, max_queue_size, workers, use_multiprocessing)
   1177 with tf.profiler.experimental.Trace(
```

```
   1178        'train',
   1179        epoch_num=epoch,
   1180        step_num=step,
   1181        batch_size=batch_size,
   1182        _r=1):
   1183   callbacks.on_train_batch_begin(step)
-> 1184   tmp_logs = self.train_function(iterator)
   1185   if data_handler.should_sync:
   1186     context.async_wait()

File ~/.local/lib/python3.8/site-packages/tensorflow/python/eager/def_function.
 ↪py:885, in Function.__call__(self, *args, **kwds)
    882 compiler = "xla" if self._jit_compile else "nonXla"
    884 with OptionalXlaContext(self._jit_compile):
--> 885   result = self._call(*args, **kwds)
    887 new_tracing_count = self.experimental_get_tracing_count()
    888 without_tracing = (tracing_count == new_tracing_count)

File ~/.local/lib/python3.8/site-packages/tensorflow/python/eager/def_function.
 ↪py:917, in Function._call(self, *args, **kwds)
    914     self._lock.release()
    915     # In this case we have created variables on the first call, so we run
 ↪the
    916     # defunned version which is guaranteed to never create variables.
--> 917     return self._stateless_fn(*args, **kwds)  # pylint:
 ↪disable=not-callable
    918   elif self._stateful_fn is not None:
    919     # Release the lock early so that multiple threads can perform the call
    920     # in parallel.
    921     self._lock.release()

File ~/.local/lib/python3.8/site-packages/tensorflow/python/eager/function.py:
 ↪3039, in Function.__call__(self, *args, **kwargs)
   3036 with self._lock:
   3037   (graph_function,
   3038    filtered_flat_args) = self._maybe_define_function(args, kwargs)
-> 3039 return graph_function._call_flat(
   3040     filtered_flat_args, captured_inputs=graph_function.captured_inputs)

File ~/.local/lib/python3.8/site-packages/tensorflow/python/eager/function.py:
 ↪1963, in ConcreteFunction._call_flat(self, args, captured_inputs,
 ↪cancellation_manager)
   1959 possible_gradient_type = gradients_util.PossibleTapeGradientTypes(args)
   1960 if (possible_gradient_type == gradients_util.POSSIBLE_GRADIENT_TYPES_NONE
   1961     and executing_eagerly):
   1962   # No tape is watching; skip to running the function.
-> 1963   return self._build_call_outputs(self._inference_function.call(
   1964       ctx, args, cancellation_manager=cancellation_manager))
```

```
      1965 forward_backward = self._select_forward_and_backward_functions(
      1966       args,
      1967       possible_gradient_type,
      1968       executing_eagerly)
      1969 forward_function, args_with_tangents = forward_backward.forward()

File ~/.local/lib/python3.8/site-packages/tensorflow/python/eager/function.py:
  ↪591, in _EagerDefinedFunction.call(self, ctx, args, cancellation_manager)
       589 with _InterpolateFunctionError(self):
       590   if cancellation_manager is None:
--> 591     outputs = execute.execute(
       592         str(self.signature.name),
       593         num_outputs=self._num_outputs,
       594         inputs=args,
       595         attrs=attrs,
       596         ctx=ctx)
       597   else:
       598     outputs = execute.execute_with_cancellation(
       599         str(self.signature.name),
       600         num_outputs=self._num_outputs,
      (…)
       603         ctx=ctx,
       604         cancellation_manager=cancellation_manager)

File ~/.local/lib/python3.8/site-packages/tensorflow/python/eager/execute.py:59⌐
  ↪in quick_execute(op_name, num_outputs, inputs, attrs, ctx, name)
        57 try:
        58   ctx.ensure_initialized()
---> 59   tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle, device_name, op_name
        60                                       inputs, attrs, num_outputs)
        61 except core._NotOkStatusException as e:
        62   if name is not None:

KeyboardInterrupt:
```

```
[23]:  # Model inference on test data
       def predict_test_data(test_dir):
           test_files = glob.glob(os.path.join(test_dir, '*.tif'))

           predictions = []
           ids = []

           for img_path in test_files[:]:
               img = load_image(img_path)
               if img is not None:
                   img = np.expand_dims(img, axis=0) / 255.0  # Normalize
                   pred = model.predict(img)
```

```
            predictions.append(pred[0][0])
        else:
            predictions.append(None)  # Handle non-loadable images
        ids.append( os.path.basename(img_path).replace('.tif', '') )

    return pd.DataFrame({
        'id': ids,
        'label': [ round(x) for x in predictions ]
    })

# Predict and save to submission file
submission_df = predict_test_data(TEST_DIR)
print(len(os.listdir(TEST_DIR)))

submission_df.to_csv(SUBMISSION_CSV, index=False)
print("Submission file saved.")
```

```
57458
Submission file saved.
```

1 Metastatic Cancer Detection Model

1.1 Model Architecture

The model is called a Convolutional Neural Network (CNN). It has several important parts:

1. **Convolutional Layers**: There are three of these layers. They look for features in the images using filters. The first layer has 32 filters, the second has 64, and the third has 128.
2. **Max Pooling Layers**: After each convolutional layer, there is a max pooling layer. This layer makes the images smaller and helps the model work faster.
3. **Flatten Layer**: This layer takes the 2D image data and turns it into a single line of numbers. This is needed for the next layers.
4. **Dense Layer**: This layer connects all the numbers together. It has 128 units and uses a special function called ReLU to help the model learn better.
5. **Output Layer**: The final layer has one unit that tells us if there is cancer or not. It gives a score between 0 and 1. If the score is close to 1, it means cancer is present.

1.2 Training Process

1. **Data Preparation**: We load the images in small groups so we don't use too much memory.
2. **Training**: We teach the model using training data. It learns to recognize signs of cancer.
3. **Validation**: We check how well the model is doing with a different set of images. This helps us see if it is learning correctly.
4. **Note on Training**: Training is done on separate environment, achieves 51% accuracy.

1.3 Inference Steps

1. **Load Test Images**: We get the images we want to test.

2. **Make Predictions**: The model looks at each test image and decides if it shows cancer or not.
3. **Save Results**: We save these predictions in a file to send out.

This model helps us find metastatic cancer in small image patches from larger scans.

[ ]: