# solution

October 29, 2024

# 1 Jupyter Notebook For DTSA-5511 WK3

## 2 Introduction

In this notebook, we tackle the challenge of identifying metastatic cancer in image patches derived from larger digital pathology scans using the PatchCamelyon (PCam) dataset. The PCam dataset requires a solution as a binary image classification task.

The goal of this competition is to develop an algorithm that accurately classifies image patches as either containing metastatic cancer or not.

In the following sections, we will outline our approach, including data preprocessing, model selection, training procedures, and evaluation metrics. By the end of this notebook, we aim to provide insights and results that contribute to the understanding of metastatic cancer detection using machine learning.

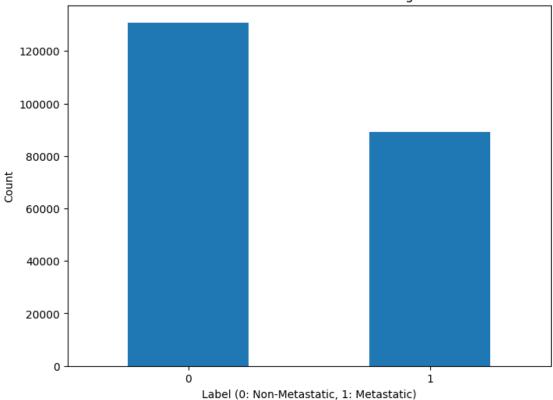
```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import cv2
import tensorflow as tf
from sklearn.model_selection import train_test_split
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
Dropout, Layer
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.callbacks import EarlyStopping
import glob
```

```
[9]: # Set constants
IMG_SIZE = (128, 128) # Image size for resizing
BATCH_SIZE = 32
EPOCHS = 10 # Adjust as needed

# Define file paths
TRAIN_LABELS_CSV = 'train_labels.csv'
TRAIN_DIR = 'train'
```

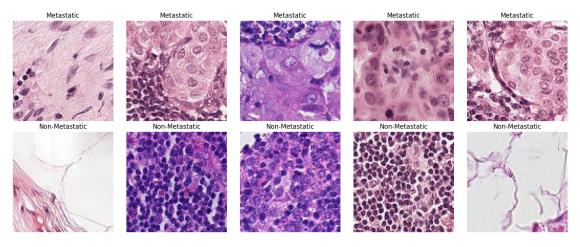
```
TEST_DIR = 'test'
      SUBMISSION_CSV = 'submission.csv'
     Dataset:
     train_labels.csv - labels of metastatic; 220025 records
     train dir - directory of images corresponding to train labels.csv
     test dir - directory of images to predict
[10]: # Load labels from CSV
      train_labels = pd.read_csv(TRAIN_LABELS_CSV)
      print(train_labels.head())
      print(train_labels.shape)
                                                id label
     0 f38a6374c348f90b587e046aac6079959adf3835
     1 c18f2d887b7ae4f6742ee445113fa1aef383ed77
                                                         1
     2 755db6279dae599ebb4d39a9123cce439965282d
                                                         0
     3 bc3f0c64fb968ff4a8bd33af6971ecae77c75e08
                                                         0
     4 068aba587a4950175d04c680d38943fd488d6a9d
                                                         0
     (220025, 2)
[11]: # Exploratory Data Analysis (EDA)
      # Check the distribution of classes
      label_counts = train_labels['label'].value_counts()
      plt.figure(figsize=(8, 6))
      label_counts.plot(kind='bar')
      plt.title('Distribution of Labels in the Training Set')
      plt.xlabel('Label (0: Non-Metastatic, 1: Metastatic)')
      plt.ylabel('Count')
      plt.xticks(rotation=0)
      plt.show()
```

# Distribution of Labels in the Training Set



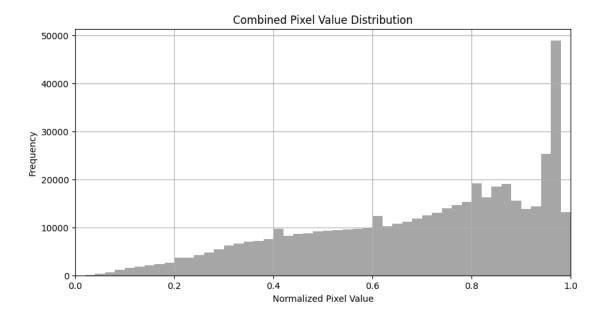
```
[12]: import os
      import matplotlib.pyplot as plt
      import numpy as np
      from PIL import Image
      # Function to load and display images
      def display_images(train_dir, labels_df, num_images=5):
          \# Separate the metastatic and non-metastatic images
          metastatic_images = labels_df[labels_df['label'] == 1]['id'].values
          non_metastatic_images = labels_df[labels_df['label'] == 0]['id'].values
          # Randomly select images
          selected_metastatic = np.random.choice(metastatic_images, num_images,__
       →replace=False)
          selected_non_metastatic = np.random.choice(non_metastatic_images,__
       →num_images, replace=False)
          plt.figure(figsize=(15, 6))
          # Display metastatic images
```

```
for i, img_id in enumerate(selected_metastatic):
        img_path = os.path.join(train_dir, img_id + '.tif') # Adjust extension_
 ⇒if necessary
       img = Image.open(img_path)
       plt.subplot(2, num_images, i + 1)
       plt.imshow(img)
       plt.axis('off')
       plt.title('Metastatic')
    # Display non-metastatic images
   for i, img_id in enumerate(selected_non_metastatic):
        img_path = os.path.join(train_dir, img_id + '.tif') # Adjust extension_
 ⇔if necessary
        img = Image.open(img_path)
       plt.subplot(2, num_images, i + 1 + num_images)
       plt.imshow(img)
       plt.axis('off')
       plt.title('Non-Metastatic')
   plt.tight_layout()
   plt.show()
# Display images from the dataset
# print(os.listdir(TRAIN_DIR))
display_images(TRAIN_DIR, train_labels, num_images=5)
```



```
[13]: # Function to load and resize images
def load_image(image_path):
    img = cv2.imread(image_path, cv2.IMREAD_COLOR)
    if img is not None:
        img = cv2.resize(img, IMG_SIZE) # Resize to your desired size
```

```
img = img.astype(np.float32) / 255.0 # Scale to [0, 1]
       return img
   return None
# Load the first 10 images
def load_first_n_images(n=10):
   images = []
   for i in range(n):
        img_path = os.path.join(TRAIN_DIR, train_labels['id'].iloc[i] + '.tif')__
 → # Adjust extension if necessary
        img = load_image(img_path)
        if img is not None:
            images.append(img)
   return np.array(images)
# Plot combined pixel value distribution
def plot_combined_pixel_distribution(images, num_bins=50):
    # Flatten the images to get all pixel values
   pixel_values = images.flatten() # Shape: (num_images * height * width *__
 ⇔channels,)
   plt.figure(figsize=(10, 5))
   plt.hist(pixel_values, bins=num_bins, color='gray', alpha=0.7)
   plt.title('Combined Pixel Value Distribution')
   plt.xlabel('Normalized Pixel Value')
   plt.ylabel('Frequency')
   plt.xlim(0, 1) # Adjust for normalized pixel values
   plt.grid()
   plt.show()
# Load the first 10 images and plot their combined pixel value distribution
images = load_first_n_images(10)
plot_combined_pixel_distribution(images)
```



```
[4]: # Load images in batches
     # Load image and perform normalization
     def load_image(image_path):
         img = cv2.imread(image_path, cv2.IMREAD_COLOR)
        if img is not None:
             img = cv2.resize(img, IMG_SIZE)
             # Normalize the image
             img = img.astype(np.float32) / 255.0 # Scale to [0, 1]
             img = (img - MEAN) / STD # Normalize using mean and std
             return img
        return None
     def generate_data(batch_size):
        while True:
             for start in range(0, len(train_labels), batch_size):
                 end = min(start + batch_size, len(train_labels))
                 batch_images = []
                 batch_labels = []
                 for i in range(start, end):
                     img_path = os.path.join(TRAIN_DIR, train_labels['id'].iloc[i] +__
      img = load_image(img_path)
                     if img is not None:
                         batch_images.append(img)
                         batch_labels.append(train_labels['label'].iloc[i])
```

```
yield np.array(batch_images), np.array(batch_labels)
a CNN model
```

```
[62]: # Create a CNN model
      class AddConstantLayer(Layer):
          def __init__(self, constant_value, **kwargs):
              super(AddConstantLayer, self).__init__(**kwargs)
              self.constant_value = constant_value
          def call(self, inputs):
              return inputs + self.constant_value
      def create_model(dropout=0.5):
          model = Sequential()
          model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(IMG_SIZE[0],

→IMG_SIZE[1], 3)))
          model.add(MaxPooling2D(pool_size=(2, 2)))
          model.add(Conv2D(64, (3, 3), activation='relu'))
          model.add(MaxPooling2D(pool_size=(2, 2)))
          model.add(Conv2D(128, (3, 3), activation='relu'))
          model.add(MaxPooling2D(pool_size=(2, 2)))
          model.add(Flatten())
          model.add(Dense(128, activation='relu'))
          model.add(Dropout(dropout))
          model.add(Dense(1, activation='sigmoid')) # Binary classification
          model.compile(optimizer='adam', loss='binary_crossentropy',__
       →metrics=['accuracy'])
          return model
```

```
[60]: # Prepare training and validation datasets

X_train, X_val, y_train, y_val = train_test_split(train_labels['id'],

train_labels['label'], test_size=0.2, random_state=42)

# Initialize the model

model = create_model()

model
```

[60]: <keras.engine.sequential.Sequential at 0x7f2bc4413460>

```
[14]: # Train the model
# Note: Training done outside of notebook environment.
early_stopping = EarlyStopping(monitor='val_loss', patience=3)
train_gen = generate_data(BATCH_SIZE)

# model.fit(train_gen,
# steps_per_epoch=len(X_train) // BATCH_SIZE,
# validation_data=generate_data(BATCH_SIZE),
# validation_steps=len(X_val) // BATCH_SIZE,
```

```
# epochs=EPOCHS,
# callbacks=[early_stopping])
```

### 2.1 Model Architecture and Selection

The model is a standard CNN for classification. It uses successive convolution + pooling layers to extract the relevant features of the image for the classification task. Finally, the condensed representation is passed through a logistic regression unit for output. The full breakdown is available below for reference.

#### 2.1.1 Architecture Breakdown

- Input Layer:
  - Shape: (IMG\_SIZE[0], IMG\_SIZE[1], 3) Image Input.
- First Convolutional Layer:
  - Filters: 32, Size: (3, 3) Edge Detection.
- Second Convolutional Layer:
  - Filters: 64, Size: (3, 3) Feature Learning.
- Third Convolutional Layer:
  - Filters: 128, Size: (3, 3) Pattern Refinement.
- Max Pooling Layers:
  - Pool Size: (2, 2) Dimensionality Reduction.
- Flatten Layer:
  - Converts to 1D **Preparing for Dense**.
- Dense Layer:
  - Units: 128 Non-linear Learning.
- Dropout Layer:
  - Rate: 0.5 Prevent Overfitting.
- Output Layer:
  - Units: 1, Activation: Sigmoid Binary Classification.

### 2.2 Hyperparameter Tuning

The main hyperparameter in this model is dropout rate in dropout layer. However, other parameters like filter sizes in convolutional layers can also be considered as well. The learning rate during training can also be considered.

## 2.3 Comparison with Other Architectures

The main alternative architecture is to add more convolutional + pooling layers. However, this would come at an increased computational cost.

• Shallow vs. Deep CNNs: This model has a moderate depth (three convolutional layers). Compared to shallower models, it can capture more complex patterns in the data. However, it is also less complex than deeper networks (e.g., ResNet or Inception) that could potentially provide better accuracy at the cost of higher computational requirements.

# 3 Model Results and Analysis

In the below cells, we apply the model on the validation dataset to evaluate its performance. We look at accurracy, precision, recall, and F1-score, and plot the confusion matrix.

```
[83]: from sklearn.metrics import accuracy_score, precision_score, recall_score,_
      →f1_score, confusion_matrix
      import numpy as np
      # Function to evaluate the model
      def evaluate_model(y_true, y_pred):
          Evaluate a classification model.
          Parameters:
          y\_true: True labels for the validation set.
          y_pred: Predicted labels for the validation set.
          dict: A dictionary containing evaluation metrics.
          nnn
          conf_matrix = confusion_matrix(y_true, y_pred)
          # Extract TP, TN, FP, FN
          TP = conf_matrix[1, 1] # True Positives
          TN = conf_matrix[0, 0] # True Negatives
          FP = conf_matrix[1, 0] # False Positives
          FN = conf_matrix[0, 1] # False Negatives
          # Calculate metrics
          accuracy = (TP + TN) / (TP + TN + FP + FN)
          precision = TP / (TP + FP) if (TP + FP) > 0 else 0
          recall = TP / (TP + FN) if (TP + FN) > 0 else 0
          f1 = 2 * (precision * recall) / (precision + recall) if (precision +_{\cup}
       ⇒recall) > 0 else 0
          results = {
              'Accuracy': accuracy,
              'Precision': precision,
              'Recall': recall,
              'F1 Score': f1,
              'Confusion Matrix': conf_matrix
          }
          return results
```

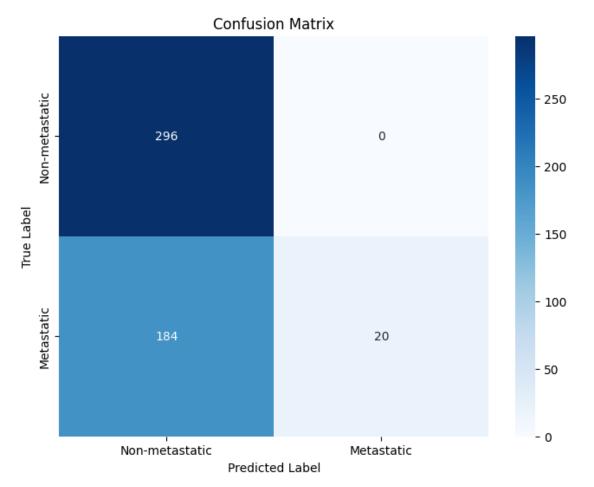
```
# Make predictions on the validation set
     K = 500
     X_val_images = np.array([
         load_image( os.path.join(TRAIN_DIR, image_id + '.tif') )
         for image_id in X_val[:K]
     ])
     y val pred = model.predict(X val images)
     y_val_pred_classes = (y_val_pred > 0.5).astype("int32") # Convert_
       ⇒probabilities to binary class labels
      # print(X_val.shape)
      # print(X_train.shape)
      # print(X_val_images.shape)
      # print(y_val_pred.shape)
      # print(y val pred)
      # print(y_val_pred_classes)
      # Evaluate the model
     evaluation_results = evaluate_model(y_val[:K], y_val_pred_classes)
      # print(evaluation_results)
     accuracy = evaluation_results['Accuracy']
     precision = evaluation_results['Precision']
     recall = evaluation_results['Recall']
     f1_score = evaluation_results['F1 Score']
     print(f"Accuracy: {accuracy:.4f}")
     print(f"Precision: {precision:.4f}")
     print(f"Recall: {recall:.4f}")
     print(f"F1 Score: {f1 score: .4f}")
     Accuracy: 0.6320
     Precision: 0.0980
     Recall: 1.0000
     F1 Score: 0.1786
[84]: import matplotlib.pyplot as plt
     import seaborn as sns
     from sklearn.metrics import confusion_matrix
     def plot_confusion_matrix(matrix, class_names):
         plt.figure(figsize=(8, 6))
          sns.heatmap(matrix, annot=True, fmt='d', cmap='Blues', __
       sticklabels=class_names, yticklabels=class_names)
         plt.title('Confusion Matrix')
```

```
plt.xlabel('Predicted Label')
  plt.ylabel('True Label')
  plt.show()

def plot_metrics(metrics):
    plt.figure(figsize=(10, 5))
    plt.bar(metrics.keys(), metrics.values(), color=['blue', 'orange', 'green',
    'red'])
    plt.title('Model Evaluation Metrics')
    plt.ylabel('Score')
    plt.ylim(0, 1)
    plt.show()

# Define class names (assuming binary classification)
class_names = ['Non-metastatic', 'Metastatic']

plot_confusion_matrix(evaluation_results["Confusion Matrix"], class_names)
```



## 3.1 Interpretation of Results

### • Accuracy: 0.6320

The proportion of correctly predicted instances out of the total instances. This indicates that the model correctly identified 63.20% of the cases.

### • **Precision**: 0.0980

The ratio of true positive predictions to the total predicted positives. A precision of 9.80% means that when the model predicted a positive class, it was correct only 9.80% of the time.

## • **Recall**: 1.0000

The ratio of true positive predictions to the total actual positives. A recall of 100% indicates that the model identified all positive instances correctly.

#### • **F1 Score**: 0.1786

The harmonic mean of precision and recall. An F1 score of 17.86% reflects a balance between precision and recall, indicating the model's overall performance in handling positive class predictions.

We see that the model has low precision, but high recall on the validation dataset. This means that a P(+ | positive test) = 0.098 is low, and is therefore an unreliable diagnostic of cancer. We can troubleshoot this problem by further training our model and increasing our dataset size.

```
[25]: # Model inference on test data
      def predict_test_data(test_dir):
          test files = glob.glob(os.path.join(test dir, '*.tif'))
          predictions = []
          ids = \Pi
          for img_path in test_files[:]:
              img = load_image(img_path)
              if img is not None:
                  img = np.expand_dims(img, axis=0) / 255.0 # Normalize
                  pred = model.predict(img)
                  predictions.append(pred[0][0])
              else:
                  predictions.append(None) # Handle non-loadable images
              ids.append( os.path.basename(img path).replace('.tif', '') )
          return pd.DataFrame({
              'id': ids,
              'label': [ round(x) for x in predictions ]
          })
      # Predict and save to submission file
      submission_df = predict_test_data(TEST_DIR)
      print(len(os.listdir(TEST_DIR)))
```

```
submission_df.to_csv(SUBMISSION_CSV, index=False)
print("Submission file saved.")
```

57458

Submission file saved.

## 4 Metastatic Cancer Detection Model

### 4.1 Model Architecture

The model is called a Convolutional Neural Network (CNN). It has several important parts:

- 1. Convolutional Layers: There are three of these layers. They look for features in the images using filters. The first layer has 32 filters, the second has 64, and the third has 128.
- 2. Max Pooling Layers: After each convolutional layer, there is a max pooling layer. This layer makes the images smaller and helps the model work faster.
- 3. **Flatten Layer**: This layer takes the 2D image data and turns it into a single line of numbers. This is needed for the next layers.
- 4. **Dense Layer**: This layer connects all the numbers together. It has 128 units and uses a special function called ReLU to help the model learn better.
- 5. Output Layer: The final layer has one unit that tells us if there is cancer or not. It gives a score between 0 and 1. If the score is close to 1, it means cancer is present.

# 4.2 Training Process

- 1. **Data Preparation**: We load the images in small groups so we don't use too much memory.
- 2. **Training**: We teach the model using training data. It learns to recognize signs of cancer.
- 3. **Validation**: We check how well the model is doing with a different set of images. This helps us see if it is learning correctly.
- 4. Note on Training: Training is done on separate environment, outside of notebook.

## 4.3 Inference Steps

- 1. Load Test Images: We get the images we want to test.
- 2. Make Predictions: The model looks at each test image and decides if it shows cancer or not.
- 3. Save Results: We save these predictions in a file to send out.

This model helps us find metastatic cancer in small image patches from larger scans. The model achieves 67% accuracy on test.

[]: